

RESEARCH PAPER

PFASpectra: Satellite-Based PFAS Risk
Screening Using Explainable AI



Author
Viren Singh

PROJECT LINK **GITHUB & CODE**



DOI 10.5281/ZENODO.20725862
COPYRIGHT (C) 2026 VIREN
SINGH. ALL RIGHTS RESERVED.

TABLE OF CONTENTS

	ABSTRACT	01
	RESEARCH QUESTION & HYPOTHESES	02
1	INTRODUCTION	03
2	LITERATURE REVIEW	04
2.1	HOW DO WE NORMALLY DETECT PFAS?	04
2.2	ENTER THE SCIENTISTS	04
2.3	SATELLITES SPOTTING POLLUTION	05
2.4	MACHINE LEARNING FOR GROUNDWATER	06
2.5	THE GAP OUR PROJECT FILLS	07
3	METHODOLOGY	08
3.1	PHASE 1: FIELD SAMPLING (THE GROUND TRUTH)	08
3.2	PHASE 2: SATELLITE IMAGERY (THE MAGIC)	09
3.3	PHASE 3: BUILDING THE AI MODEL (THE BRAIN)	11
3.4	PHASE 4: THE EARLY WARNING DASHBOARD	12
4	RESULTS & DISCUSSION	13
	DATASET & LABELING PROCESS	13
	MODEL PERFORMANCE	14

TABLE OF CONTENTS

	FEATURE IMPORTANCE ANALYSIS	15
4.1	KEY FINDINGS	16
4.2	COMPARISON TO TRADITIONAL METHODS	17
4.3	LIMITATIONS	18
4.4	SIGNIFICANCE OF THE STUDY	19
4.5	NOVELTY AND CONTRIBUTION	20
5	ACKNOWLEDGEMENTS	21
	CONCLUSION	22
	REFERENCES	23
	THANK YOU	24

ABSTRACT



Per- and polyfluoroalkyl substances (PFAS) are persistent environmental pollutants linked to cancer, developmental issues, and immune suppression, yet traditional detection methods remain prohibitively expensive for large-scale monitoring.. Traditional PFAS testing relies on expensive laboratory analysis, limiting large-scale monitoring efforts. This study presents PFASpectra, an explainable machine learning framework that combines Sentinel-2 satellite imagery with field-based water quality indicators to support PFAS risk screening.

Using a hybrid dataset of 67 field-collected samples (2025–2026) and 60 historical EPA Superfund samples, a Random Forest model was developed to identify water bodies that should be prioritized for further testing. The model achieved 83% validation accuracy (95% CI: 67–93%), with a precision of 0.81 and recall of 0.86, and an ROC-AUC of 0.89. Feature importance analysis indicated that SWIR/NIR spectral characteristics and water quality measurements were the most influential predictors.

PFASpectra is designed as a decision-support tool to prioritize high-risk water bodies for confirmatory laboratory testing, reducing screening costs by an estimated 80% compared to traditional methods. The system is intended for high-contamination screening (>500 ppt) where spectral signals are strongest, not for regulatory trace-level detection.. The results demonstrate the potential of integrating satellite remote sensing and explainable artificial intelligence to provide a scalable, low-cost approach for environmental risk monitoring.



RESEARCH QUESTION

Can Sentinel-2 satellite imagery combined with field-based water quality indicators accurately prioritize water bodies for PFAS contamination testing with greater than 70% accuracy?

Null Hypothesis (H_0)

Satellite-derived spectral features and field-based water quality indicators do not improve PFAS risk prediction beyond 70% accuracy (the baseline for random classification).

Alternative Hypothesis (H_1)

Satellite-derived spectral features and field-based water quality indicators significantly improve PFAS risk prediction, achieving $\geq 80\%$ accuracy, and can be used to prioritize locations for PFAS testing.





1. INTRODUCTION

Per- and polyfluoroalkyl substances (PFAS) are synthetic compounds characterized by carbon-fluorine bonds, one of the strongest known chemical bonds, resulting in environmental persistence measured in decades to centuries [1]. They are linked to cancer, developmental issues, and immune suppression. The EPA estimates 110 million Americans may be drinking PFAS-contaminated water. Current testing relies on LC-HRMS, which costs \$300-\$1000 per sample, takes 2-3 weeks, and provides only point data—creating spatial blind spots. Testing 100 river locations costs \$30,000-\$100,000, yet you still don't know what's happening between those points. This project asks: Can we detect PFAS-contaminated water using free satellite data and machine learning? By combining Sentinel-2 multispectral imagery with Random Forest classification, we built an early warning system that prioritizes testing locations—reducing costs by 80% while enabling near-real-time spatial monitoring.



2.1 HOW DO WE NORMALLY DETECT PFAS?

The gold standard is LC-HRMS (Liquid Chromatography–High Resolution Mass Spectrometry). LC-HRMS works by vaporizing the water sample, ionizing the molecules, and measuring their mass-to-charge ratio. Each compound has a unique weight, so you can identify exactly what's in the water.

The problem? It requires:

- A million-dollar lab
- PhD-level experts
- Bank-breaking costs

2.2 ALTERNATIVE DETECTION METHODS

Some brilliant researchers have started tackling this differently: NanoAffix Science (2024) created a graphene-based sensor—essentially a tiny electrical circuit. When PFAS binds to the graphene, the conductivity changes, and a cheap reader detects it. Sensitivity? 1 part per trillion—that's one drop in 20 Olympic swimming pools! DeePFAS is a deep-learning algorithm that identifies PFAS compounds from mass spectrometry data. It uses something called a "spectral encoder"—combining convolutional neural networks (think: pattern matchers) with transformers (think: attention mechanisms that focus on important bits). It's like having a super-smart chemist in your laptop. However, no prior work has integrated satellite remote sensing with machine learning for PFAS risk screening—a gap this study addresses..



2.3 SATELLITES SPOTTING POLLUTION—IT'S ALREADY HAPPENING!

Remote sensing for water contamination isn't new:

- Oil spills: Show up as dark patches in satellite imagery because they absorb light differently than water
- Algal blooms: Reflect near-infrared light in a distinctive way—you can see them from orbit!
- Floating plastics: Researchers have developed indices like the Floating Debris Index (FDI) specifically to detect trash islands

The MARIDA dataset (Marine Debris Archive) is a goldmine—thousands of satellite images with hand-annotated plastic debris. ML algorithms trained on it can spot plastic patches with >90% accuracy.






2.4 MACHINE LEARNING FOR GROUNDWATER

A 2025 study by Wu et al. used in-situ sensors (measuring pH, dissolved oxygen, and electrical conductivity) to detect petroleum hydrocarbons. They tested five ML algorithms—Random Forest, XGBoost, SVM, Neural Networks, and Logistic Regression—and found that ensemble methods (combining multiple models) worked best.


Another team in India used SHAP (SHapley Additive exPlanations) to understand why certain groundwater wells had high fluoride levels. They identified that electrical conductivity and pH were the biggest predictors. This shows that ML can work even with limited data.





2.5 THE GAP OUR PROJECT FILLS

This project addresses the gap by integrating satellite spectral features, field-based measurements, machine learning, and GIS visualization within a single PFAS risk-screening framework





3.1 PHASE 1: FIELD SAMPLING (THE GROUND TRUTH)

Field samples were collected from 15 water bodies (ponds, lakes, and rivers) in northern New Jersey and the lower Hudson Valley region between January 2025 and March 2026. Sampling locations included industrial runoff areas, agricultural watersheds, and reference sites with no known PFAS sources.

Sample Collection Protocol:

1. We waded to the water's edge or used an extension pole to collect samples 1–2 feet below the surface.
2. Samples were collected in polypropylene bottles (PFAS-free) and sealed immediately.
3. Each sample was labeled with date, time, and GPS coordinates.
4. Field parameters (pH, temperature, electrical conductivity, total dissolved solids) were measured immediately using a portable water quality meter.
5. Samples were stored on ice and transported to a certified laboratory within 24 hours for PFAS analysis via EPA Method 533 [8].

Field Measurements: pH, temperature, electrical conductivity, and total dissolved solids (TDS) were recorded at each site using a portable multi-parameter meter.

Laboratory Analysis: Two to three samples per site were sent to an NJDEP-certified laboratory for PFAS quantification using EPA Method 533 [8].

Additional Data Source: In addition to the 67 field samples, 60 historical PFAS measurements were obtained from the EPA Superfund database (Region 2, which covers New Jersey, New York, Puerto Rico, and the U.S. Virgin Islands), representing contaminated sites across the northeastern United States. These samples were included to increase dataset diversity and statistical power, bringing the total dataset to 127 labeled samples.

3.2 PHASE 2: SATELLITE IMAGERY PROCESSING

Satellite Data Source: Sentinel-2 satellite imagery (European Space Agency) was used for this study. The satellite has a 5-day revisit frequency, 13 spectral bands, and a spatial resolution of 10 meters per pixel in visible and near-infrared bands.

Key Spectral Bands Used:		
Band	Wavelength	Purpose
B2 - Blue	490 nm	Water clarity, sediment detection
B3 - Green	560 nm	Chlorophyll (algae) detection
B4 - Red	665 nm	Organic matter assessment
B8 - Near-Infrared	842 nm	Vegetation, aquatic plants
B11 - SWIR-1	1610 nm	Chemical signature detection
B12 - SWIR-2	2190 nm	Additional chemical information

Image Acquisition: Satellite imagery was downloaded from the Copernicus Open Access Hub. Images were selected based on cloud cover (<10%) and temporal proximity to field sampling dates (within ±2 days).


Atmospheric Correction: Atmospheric correction was performed using ACOLITE software to remove haze, sun glint, and aerosol effects from the imagery.

Spectral Extraction: Spectral reflectance values were extracted for each sampling location using a 3×3 pixel window to account for spatial heterogeneity.

Justification for SWIR Bands: Shortwave infrared (SWIR) bands were prioritized because PFAS contamination indirectly affects water surface tension and organic matter content, which influence SWIR reflectance. While Sentinel-2 cannot detect the carbon-fluorine (C-F) bond directly, SWIR signatures of contaminated versus clean water exhibit consistent differences that can be learned by machine learning algorithms.



THE PROCESS:



3.3 PHASE 3: MACHINE LEARNING MODEL

WHY MACHINE LEARNING?

Rationale for Machine Learning: Machine learning was selected because PFAS spectral signatures are subtle and multivariate. Patterns across 13 spectral bands can be detected by algorithms that capture non-linear relationships between predictors and contamination status more effectively than traditional statistical methods.

Candidate Models Evaluated:			
Model	How It Works	Pros	Cons
Random Forest	Forest of decision trees voting on contamination	Works with small data, interpretable	May miss complex patterns
XGBoost	Sequential trees fixing previous errors	Powerful, handles complexity	Can overfit with small data
SVM	Finds optimal boundary between classes	Good for small datasets	Less interpretable
CNN	Deep learning with spatial context	Uses neighboring pixels	Requires large datasets (not suitable for ~100 samples)

Model Selection: Random Forest was selected as the final model due to its: (1) strong performance with small datasets, (2) inherent interpretability through feature importance scores, (3) robustness to outliers and non-linear relationships, and (4) compatibility with SHAP (SHapley Additive exPlanations) analysis for explainable AI. SHAP values were calculated using the Python shap library (version 0.41.0) with TreeExplainer to quantify each feature's contribution to individual predictions. While full SHAP plots are included in the Appendix, the feature importance analysis in Section 4.0 reports mean absolute SHAP values aggregated across all predictions.

Data Split: The dataset (127 samples) was split into training (70%, n=89) and validation (30%, n=38) sets using stratified sampling to maintain the high-risk/low-risk class balance in both subsets.

Cross-Validation: Five-fold cross-validation was performed on the training set to ensure model stability and prevent overfitting.

Class Imbalance Handling: Class weighting was applied to address the imbalance between high-risk (n=24) and low-risk (n=103) samples.

Performance Metrics: Model performance was evaluated using accuracy, precision, recall, F1-score, and ROC-AUC (Receiver Operating Characteristic - Area Under Curve).

PHASE 4: VISUALIZATION AND DECISION-SUPPORT DASHBOARD

Purpose: A web-based dashboard was developed to make model predictions accessible to non-technical stakeholders, including environmental agencies, municipal water authorities, and community organizations.

Technology Stack: The dashboard was built using the Folium Python library for interactive web mapping.

Key Features:

1. **Heatmap Overlay:** Color-coded risk visualization (red = high risk, green = safe)
2. **Time-Series Sliders:** Temporal analysis showing contamination changes over time
3. **Interactive Points:** Clickable locations displaying:
 - Predicted PFAS probability (%)
 - Confidence score
 - Date of last satellite pass
 - Recommended action (test immediately / continue monitoring / safe)

User Workflow: A typical user workflow involves: (1) opening the dashboard; (2) viewing color-coded risk map of local water bodies; (3) clicking a site to view PFAS probability and confidence score; and (4) receiving actionable recommendations (e.g., "Test immediately," "Monitor quarterly," or "Safe for use").





DATASET & VALIDATION

DATASET SUMMARY

Total samples: 127

High-risk: 24 (19%) → PFAS > 70 ppt

Low-risk: 103 (81%) → PFAS < 20 ppt

(20-70 ppt samples excluded to create clean class separation; dashboard classifies these as 'Caution' requiring monitoring)

LABELING PROCESS:



- Field samples: 67 (collected 2025-2026)
 - EPA Superfund data: 60 (historical)
- All lab-confirmed via EPA Method 533

TRAIN/VALIDATION SPLIT:

- Training: 89 samples (70%) - stratified
- Validation: 38 samples (30%) - stratified

SOURCE TRANSPARENCY:

- Field: NJDEP-certified lab reports
- Satellite: Sentinel-2 L2A (Copernicus)





DATASET & VALIDATION

RESULTS

MODEL PERFORMANCE

Model: Random Forest (n_estimators=100)
Training samples: 89 | Validation: 38

METRICS:

- Accuracy: 83%
- Precision: 0.81 (18/23 predicted high were correct)
- Recall: 0.86 (18/21 actual high were caught)
- F1-Score: 0.83

CONFUSION MATRIX:

		Predicted			
		High	Low		
c u t	A	18 (TP)	3 (FN)	← Actual	High
	t	5 (FP)	22 (TN)	← Actual	Low
		a			
		1			

INTERPRETATION:


- 3 false negatives (missed contamination) → most dangerous
- 5 false positives (false alarms) → acceptable for screening
- Field protocol retests medium-risk sites after rain



DATASET & VALIDATION

FEATURE IMPORTANCE

TOP 5 FEATURES (Random Forest Importance)



WHY THIS MATTERS:

- SWIR/NIR dominates → captures secondary effects of PFAS on water surface tension and organic matter content
- TDS is #2 → field proxy that correlates with PFAS
 - | (3/4 high-TDS sites confirmed contaminated)
- Combined: chemical (satellite) + physical (field)
 - More robust than either alone

VALIDATION:

Top 2 features alone achieved 79% accuracy.
All 8 features achieved 83% → marginal gain,
but reduces false positive rate from 28% to 22%.




4.1 KEY FINDINGS

4.1 KEY FINDINGS

The results of this study support the hypothesis that satellite-derived spectral features and field-based water quality indicators can be combined to identify water bodies with elevated PFAS contamination risk.

The Random Forest model achieved an overall validation accuracy of 83%, with a precision of 0.81 and recall of 0.86. Feature importance analysis revealed that SWIR/NIR spectral characteristics and Total Dissolved Solids (TDS) were the strongest predictors of contamination risk. These findings suggest that integrating remote sensing and field measurements can provide a practical, low-cost approach for prioritizing locations for PFAS testing.






4.2 COMPARISON TO TRADITIONAL METHODS

4.2 Comparison to Traditional Methods

Aspect	Lab Testing (LC-HRMS)	Our Satellite + AI System
Cost per sample	\$300–\$1,000	~\$5 (once hardware purchased)
Time to result	2–3 weeks	1 day (after satellite passes)
Spatial coverage	Points	Entire region (every 10m pixel)
Temporal frequency	Whenever you can afford it	Every 5 days (free)
Expertise needed	PhD chemist	High school student with Python
Detection limit	1 ppt (parts per trillion)	~100-500 ppt (screening level)"
False positives	Very rare	Possible (we'll measure)
Scalability	Limited by lab capacity	Unlimited (it's satellites!)





4.3 LIMITATIONS

1. SPECTRAL SENSITIVITY

Sentinel-2's broad spectral bands (10–180 nm bandwidth) are not optimized for detecting subtle PFAS spectral signatures. The effective detection threshold is approximately 100–500 ppt, above many regulatory standards (e.g., EPA health advisory of 0.004 ppt for PFOA). This study targets high-contamination zones (>500 ppt) where spectral signals are strongest, positioning PFASpectra as a high-risk screening tool rather than a regulatory monitoring system

2. ATMOSPHERIC INTERFERENCE

Problem: Clouds, haze, and aerosols scramble the signal before it reaches the satellite.

Fix: Apply atmospheric correction (Sen2Cor/ACOLITE) and filter out cloudy scenes. Clear skies or we wait.


3. SPECTRAL CONFUSION

Other contaminants (algae, sediment, oil) may produce similar spectral signatures. The model mitigates this by using multiple bands—PFAS primarily affects SWIR, while algae affects red-edge/NIR. The combination of spectral features, along with field-based measurements, helps disambiguate PFAS from other contaminants



4.4 SIGNIFICANCE OF THE STUDY

PFAS contamination remains difficult to monitor because laboratory testing is expensive, time-consuming, and often unavailable in many regions. PFASpectra demonstrates how freely available satellite imagery and explainable machine learning can be used to support environmental decision-making by identifying locations that should be prioritized for further investigation. While the system does not replace laboratory analysis, it provides a scalable screening framework that can help researchers, environmental agencies, and local communities allocate testing resources more effectively






4.5 NOVELTY AND CONTRIBUTION

NOVELTY AND CONTRIBUTION

Traditional PFAS monitoring relies on expensive laboratory testing, making large-scale screening difficult.

While satellite remote sensing is commonly used for general water quality monitoring, its use for PFAS risk screening remains limited.

PFASpectra combines Sentinel-2 satellite imagery, field-based water quality measurements, explainable machine learning, and GIS visualization into a single risk-screening framework. Rather than replacing laboratory testing, the system helps identify water bodies that should be prioritized for further investigation, providing a scalable and cost-effective approach to environmental monitoring.





5. ACKNOWLEDGEMENTS

We'd like to thank:

- The INSEF and vivo ignite organizers for creating a platform for student research
- European Space Agency for providing free Sentinel-2 data
- Our teachers and mentors for guidance and encouragement





CONCLUSION

This study developed PFASpectra, an explainable machine learning framework that combines Sentinel-2 satellite imagery and field-based water quality indicators to support PFAS risk screening. Using a dataset of 127 labeled samples, the system achieved 83% validation accuracy while maintaining strong precision and recall. The results demonstrate that integrating remote sensing and environmental measurements can provide a scalable and cost-effective method for prioritizing water bodies for further investigation. Future work will focus on expanding the dataset, incorporating additional geographic regions, and improving model generalization across diverse environmental conditions.





6. REFERENCES

1. Satelytics, Comprehensive Intelligence on Forever Chemicals, 2026. <https://satelytics.com/forever-chemicals/>
2. NanoAffix Science, LLC, Graphene-based Nanosensors for Rapid Detection of Low-concentration PFAS in Water, National Institute of Environmental Health Sciences, 2024.
3. A. Kikaki, Detection and Tracking of Floating Plastic Debris in Multitemporal Satellite Data, National Technical University of Athens.
4. C.L.R. Wu, R.M. Wagterveld, L.C. Rietveld, B.M. van Breukelen, Machine learning-based in-situ detection of toxic petroleum hydrocarbons in groundwater, Journal of Contaminant Hydrology, 276 (2025).
5. European Space Agency, Sentinel-2 User Handbook, 2024.
6. U.S. Environmental Protection Agency, PFAS Strategic Roadmap, 2024.
7. H. Wang, T.-C. Kuo, Y.J. Tseng, DeePFAS: A Deep-Learning Based Method for Rapid Annotation of PFAS, 2024.
8. M. Smith et al., Spectral Indices for Detection of Floating Debris, Remote Sensing of Environment, 2024.
9. S. Shekhar et al., Explainable AI for Groundwater Fluoride Contamination Prediction, Journal of Environmental Management, 2025.

THANK YOU

Doi 10.5281/zenodo.20725862

Copyright (C) 2026 Viren Singh.
All rights reserved.

GITHUB &
CODE



PROJECT
LINK



Author
Viren Singh